

Spatial SQL II

CASA0025: Building Spatial Applications
with Big Data



Ollie Ballinger

Grain Theft in Ukraine

- "They take grain to the annexed Crimea first, where they transport it to Kerch or Sevastopol [ports], then they load Ukrainian grain on Russian ships and go to the **Kerch Strait**. There, in the **Kerch Strait** [between Crimea and Russia], they transfer Ukrainian grain from small ships on to bulk carriers, where it is mixed with grain from Russia - or in some cases, they sail to this area just to give the appearance they are loading up with Russian grain."

Key transport routes into Crimea and Russia



Source: Institute for the Study of War (22 June) / BBC Research

BBC

Step 1: Load Grain

- Smaller feeder vessels turn off AIS and sail to the Avlita Grain Terminal in Sevastopol.
- Using optical and radar satellite imagery, we identified **180 ships with AIS turned off** at this terminal in the first year of the war.



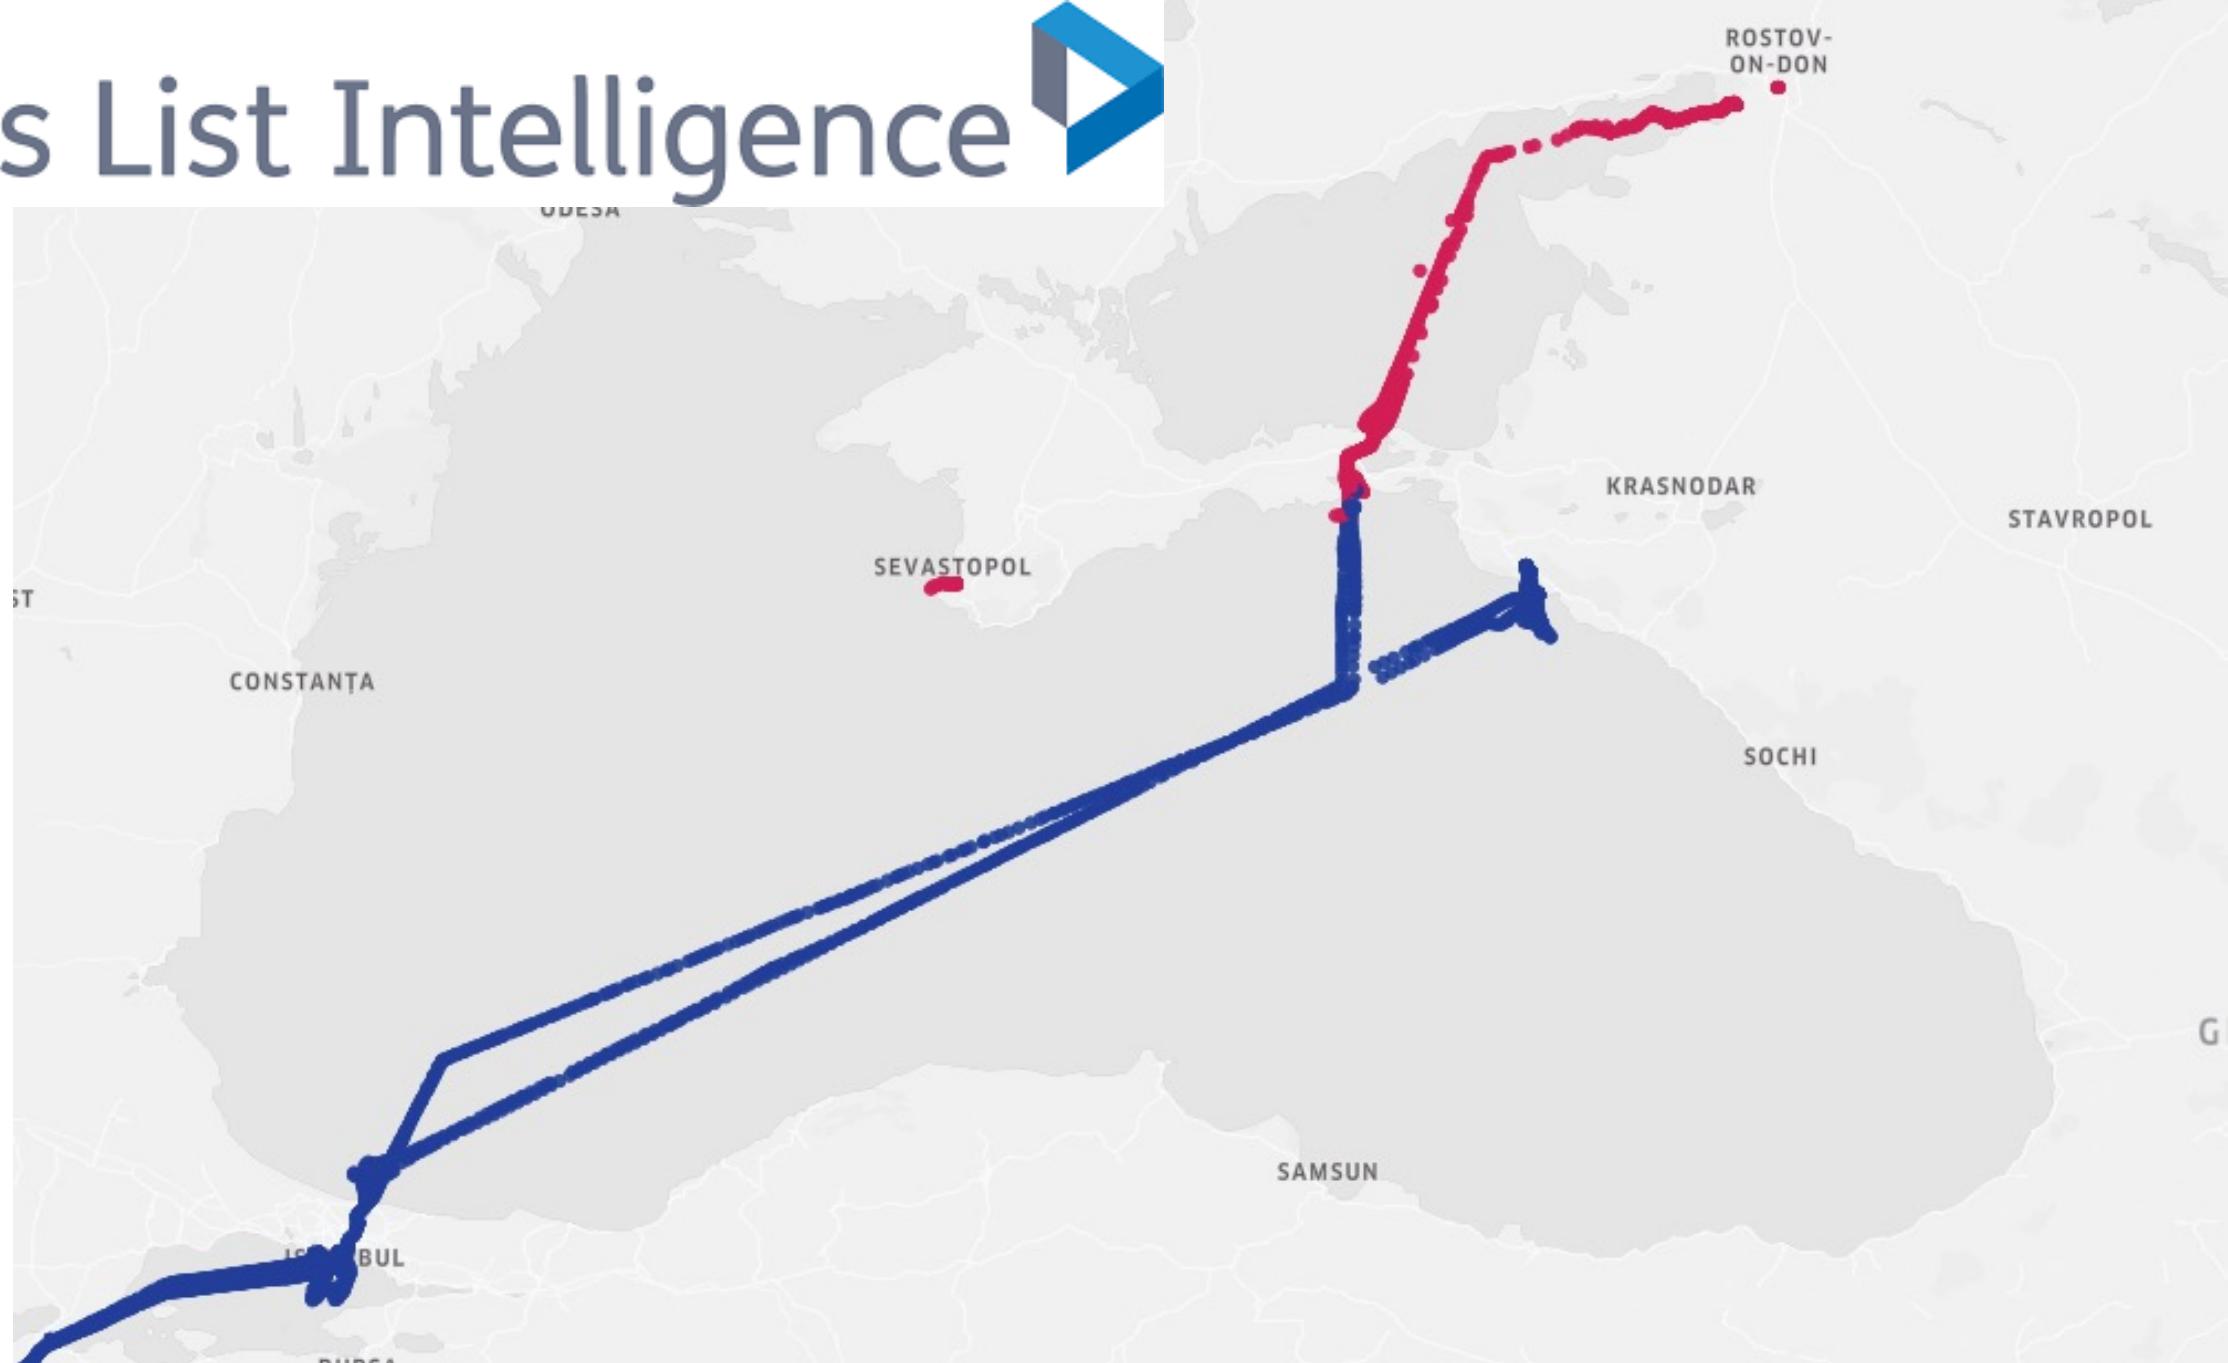
Step 2: Ship-to-Ship Transfer







Lloyd's List Intelligence



LLOYD's LIST. N° 560

FRIDAY, January 2. 1740.

THIS List, which was formerly publish'd once a Week, will now continue to be publish'd every Tuesday and Friday, with the Addition of the Stocks Course of Exchange, &c.—Subscriptions are taken in at Three Shillings per Quarter, at the Bar of *Lloyd's Coffee-House* in *Lombard-Street*.

Such Gentlemen as are willing to encourage this Undertaking, shall have them carefully deliver'd according to their Directions.

London Exchanges on

Amst. 34 11 a 10
Ditto Sight 34 7 $\frac{1}{2}$ a 8
Rott. 35 2 1
Antw. 35 11 a 36
Hamb: 33 10 2 U a 11 2 $\frac{1}{2}$

Paris — 3 $\frac{2}{4}$

Ditto at 2 U 3 $\frac{2}{4}$

Bourdeaux 3 $\frac{2}{4}$

Usance 3 $\frac{2}{4}$

Cadiz 4 $\frac{1}{4}$

Madrid 4 $\frac{1}{4}$

Bilboa 4 $\frac{1}{4}$

Leghorn 5 $\frac{1}{4}$

Genba 55

Venice 5 $\frac{1}{4}$

Lisbon 5 4 $\frac{7}{8}$ a 5

Oporto 5 4 $\frac{1}{2}$

Dublin 8

Aids in the Exchequer

18th 2 Shilling	1739	1000000	926800
18th 4 Ditto	1740	2000000	482600
Malt —	1739	750000	501014
Salt —	1734	1000000	910500

Gold in Coin —

Ditto in Barrs —

£ Pillar large —

£ Ditto Small —

£ Mexico large —

£ Ditto Small —

Silver in Barrs —

Annuitie^s

14l. per Cent at 22 $\frac{1}{2}$ Years Purchase

1704 to 1708 Inclusive 24 $\frac{1}{2}$ ditto

3 $\frac{1}{2}$ per Cent. 1 per Cent. præm.

3 per Cent. 5 $\frac{1}{4}$ Disc.

Lottery 1710.

Prizes for 3 Years from Michaelmas last are in course of Payment
Blanks for 3 Years from Michaelmas last 1l. 10s per Set.

Price of Stocks.—	Wednesday	Thursday	Friday
Bank Stock -----	138 $\frac{1}{2}$ a $\frac{3}{4}$		138 $\frac{1}{2}$
East India-----		156	156a56 $\frac{1}{4}$
South Sea -----	98 $\frac{1}{4}$		98 $\frac{1}{4}$
Ditto Anuity Old	110 $\frac{1}{2}$ a 10	110 $\frac{1}{2}$	110 $\frac{1}{2}$
Ditto — New	110 $\frac{1}{2}$ a $\frac{3}{4}$	110 $\frac{1}{4}$	110 $\frac{1}{4}$
3 per Cent. 1726			99 $\frac{1}{4}$
Annuity - 1731			
Million Bank ---	113	113	113
Equivalent -----	112	112	112
R. Ass. 100l paid in			
L: Ass. 13l paid in	10 $\frac{1}{4}$	10 $\frac{1}{4}$	10 $\frac{1}{4}$
7 per Cent E.m. Loan	98	98	98
5 per Cent. Ditto	74 $\frac{1}{4}$	74 $\frac{1}{4}$	75
Bank Circulation	21 10s od	21 10s od	21 10s od
Lottery Tickets	51 16s od	51 17s od	61 00s od

India Transfer Books open the 19th of January

Royal Assurance the 20th of January

South Sea New Annuity the 22d of January, 3 per Cent Annuities the 21st and 22d of January

South Sea Stock the 4th of February

The 5 per Cent Emperor's Loan, sells as above without the six Months Interest of a and a quarter per Cent, and 5 per Cent. part of the Principal to be paid of both, are now paying at the Bank

The India Dividend will be paid the 29th of January, South Sea New Annuities the 29th ditto, and the S. Sea Stock the 6th and 7th of February, Navy and Victualling Bills to the 30th June last are in course of Payment.

Interest per Cent	Wednesday	Thursday	Friday
3 India Bonds new	79	80	80
4 Salt Tallies	7 a $\frac{1}{2}$	7 a $\frac{1}{2}$	7 a $\frac{1}{2}$

Shill:
Præm:

San Damian + Risk Rating: ● LLI NO 349789 IMO 9274331 Flag Syria Status ● Live AIS Destination AQABA AIS ETA 29 Oct 2023 SANCTIONED

⚠ SANCTIONED | + Watchlist | Map

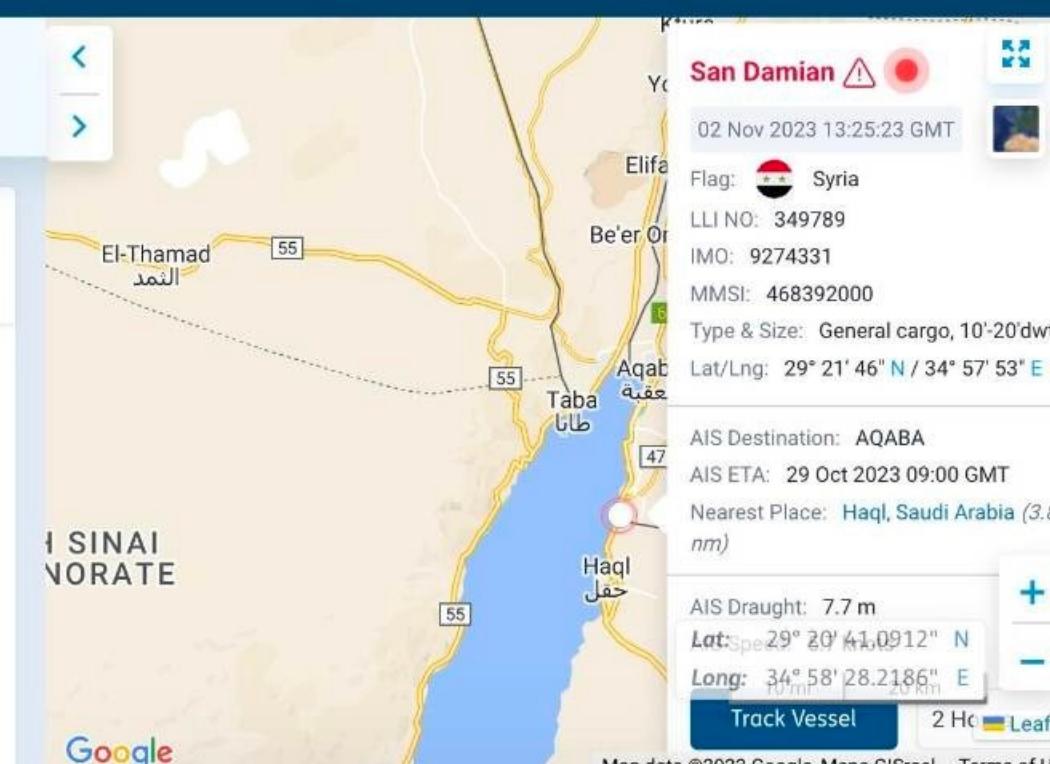
[Overview](#) [Movements](#) [Incidents](#) [Fixtures](#) [Sanctions](#) [Hull Risk](#) [Ownership](#) [STS Pairings](#)

[Summary](#) [Compliance Risks](#) [Latest News](#) [Operational Analytics](#) [Characteristics timeline](#) [Vessel Images](#) [Registration](#) [Tonnage](#)
[History](#) [Dimensions](#) [Class and Insurance](#) [Surveys](#) [Hull Details](#) [Facilities](#) [Machinery](#) [Inmarsat](#)

Summary

Last Updated: 30 Oct 2023

LLI NO:	349789	IMO:	9274331
Flag:	 Syria	LLI Vessel Type:	general cargo with container capacity
Status:	● Live	Reg. Owner:	Unknown Owners
Built:	2004	DWT:	12717
GT:	9611	Hull Type:	Single
Latest AIS message type:	A	TEU Capacity:	698



Map showing the vessel's route from El-Thamad (التمدا) to Haql (حفل), Saudi Arabia. The map includes labels for El-Thamad, Elifa, Be'er Ofer, Aqab, Taba, and Haql. The route is marked with a dashed line and road numbers 55 and 47. The vessel is currently positioned near Haql.

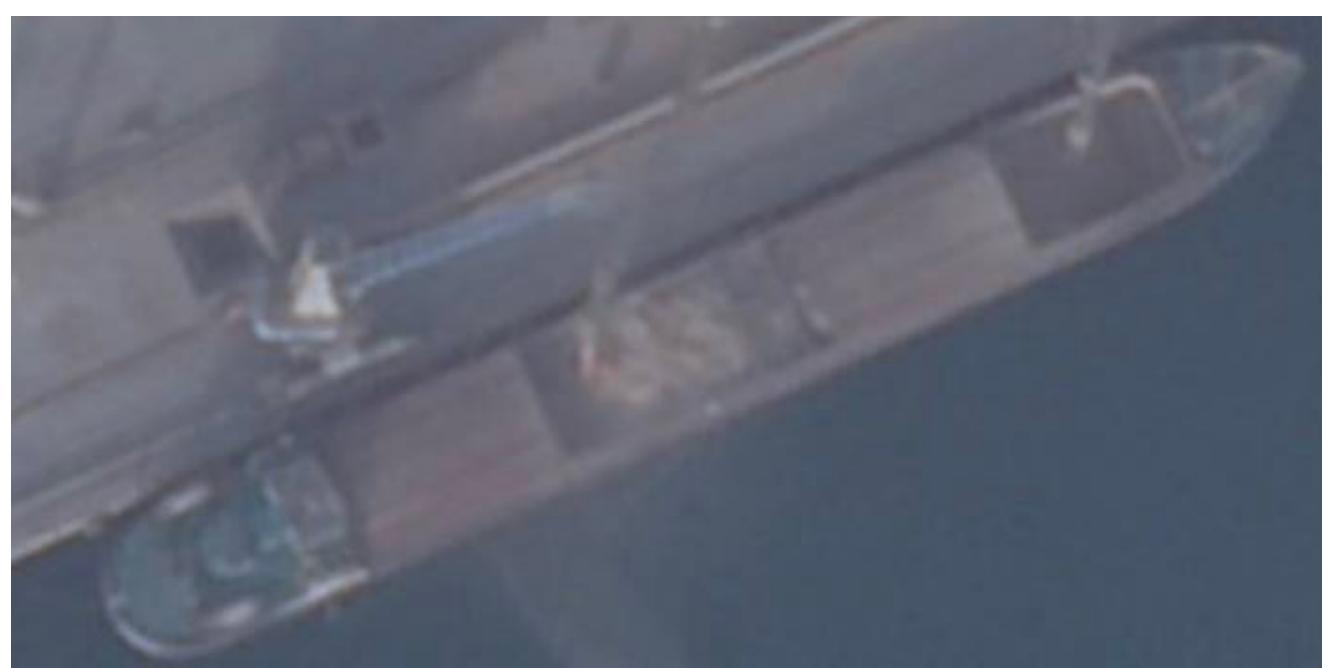
San Damian ⚠
 02 Nov 2023 13:25:23 GMT
 Flag:  Syria
 LLI NO: 349789
 IMO: 9274331
 MMSI: 468392000
 Type & Size: General cargo, 10'-20'dwt
 Lat/Lng: 29° 21' 46" N / 34° 57' 53" E

 AIS Destination: AQABA
 AIS ETA: 29 Oct 2023 09:00 GMT
 Nearest Place: Haql, Saudi Arabia (3.82 nm)

 AIS Draught: 7.7 m
 Lat: 29° 20' 41.0912" N
 Long: 34° 58' 28.2186" E

 Track Vessel 2 Hr Leaflet

Map data ©2023 Google, Mapa GIsrael Terms of Use

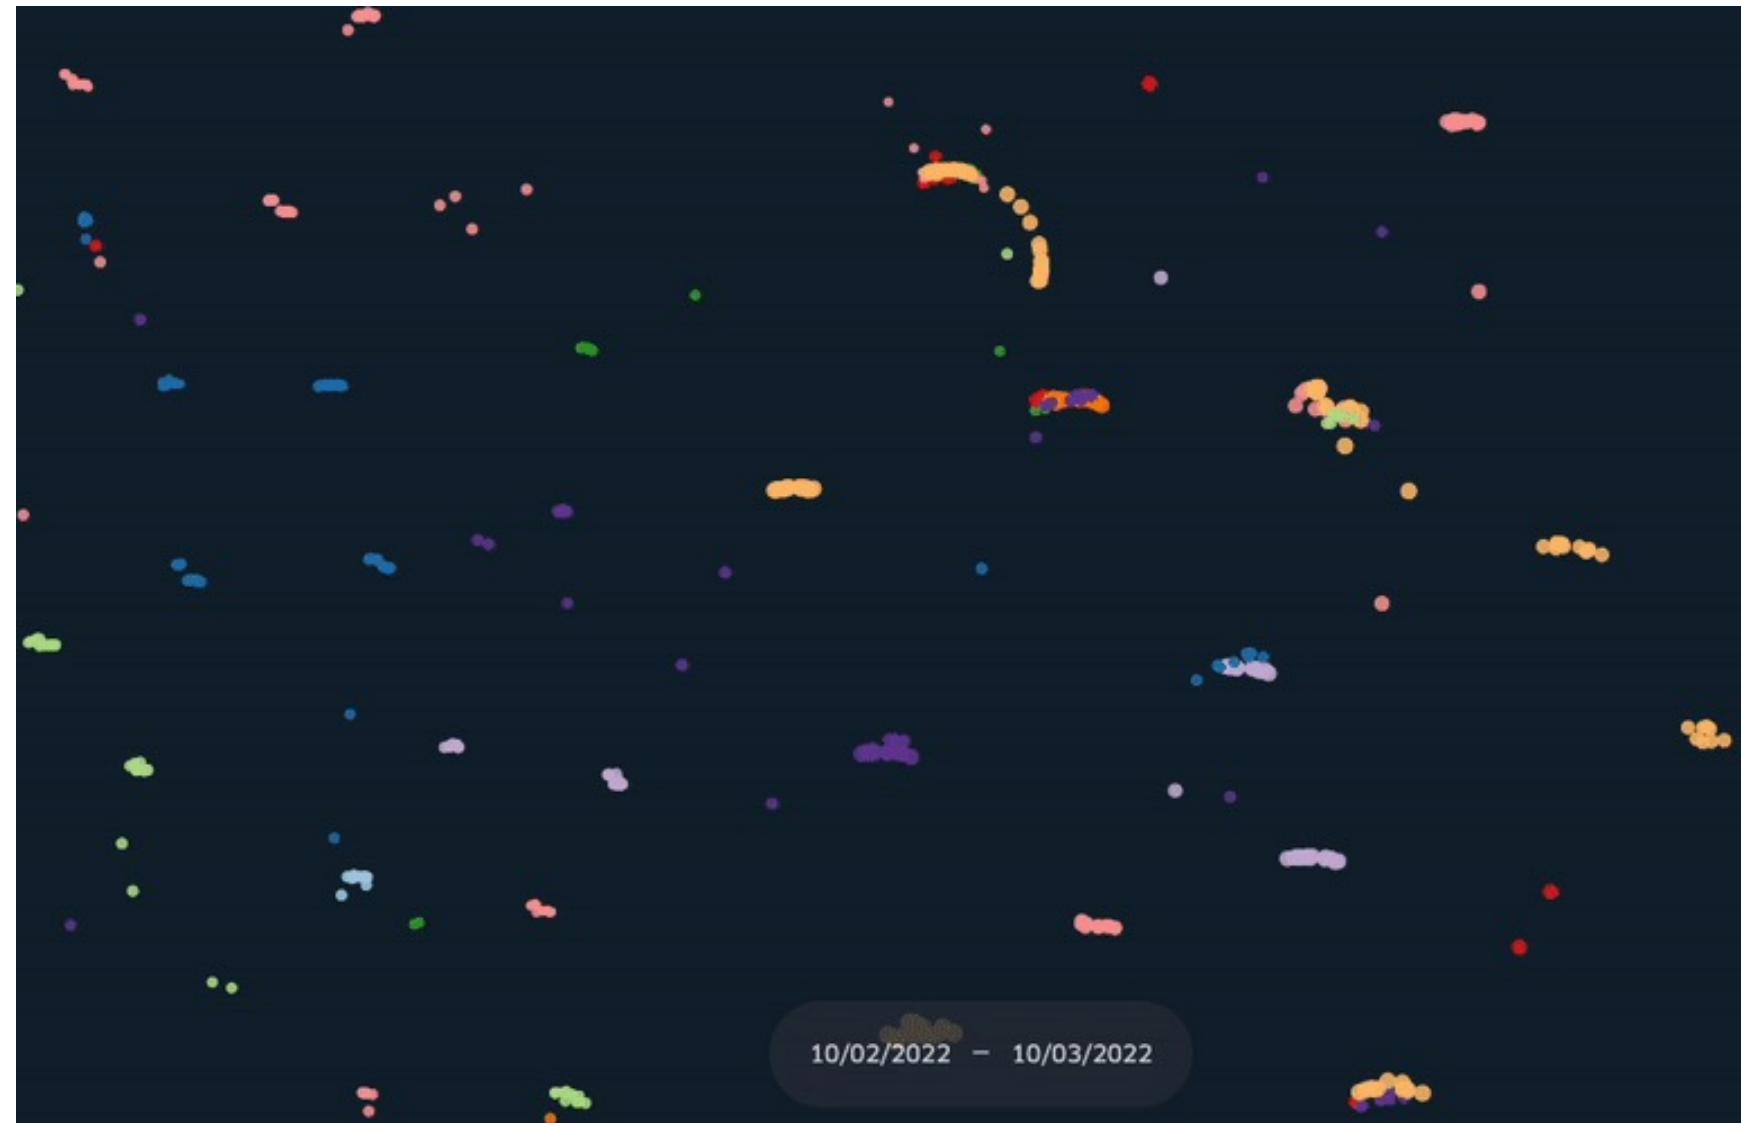


imo	vessel_name	vsl_descr	dwt	v_length	vesselid	draught	date	geometry
9372846.0	15 Temmuz	general cargo	7827.0	113.0	378595	4.4	2021-09-12 01:29:54	POINT (4063634.892855323 5636163.296415932)
9372846.0	15 Temmuz	general cargo	7827.0	113.0	378595	4.4	2021-09-12 01:36:43	POINT (4063613.00002214 5638822.792206503)
9372846.0	15 Temmuz	general cargo	7827.0	113.0	378595	4.4	2021-09-12 01:43:43	POINT (4062958.626948753 5641425.484822003)
9372846.0	15 Temmuz	general cargo	7827.0	113.0	378595	6.9	2021-09-17 00:38:34	POINT (4062702.035522474 5638123.540483131)
8956188.0	30 Let Pobedy	general cargo	5150.0	NaN	350053	3.4	2021-06-05 02:26:10	POINT (4064299.099150392 5643560.996257269)
...
9232151.0	Zeynep C	bulk carrier	53806.0	183.0	312965	6.2	2023-08-16 13:32:48	POINT (4071296.271647753 5646968.081009985)
9232151.0	Zeynep C	bulk carrier	53806.0	183.0	312965	6.2	2023-08-16 13:35:52	POINT (4070342.263611655 5643114.654042991)
9232151.0	Zeynep C	bulk carrier	53806.0	183.0	312965	6.2	2023-08-16 13:36:03	POINT (4070341.891804556 5643113.601468305)
9232151.0	Zeynep C	bulk carrier	53806.0	183.0	312965	6.2	2023-08-16 14:39:18	POINT (4071619.468864959 5645507.330939625)
9232151.0	Zeynep C	bulk carrier	53806.0	183.0	312965	6.2	2023-08-19 11:36:03	POINT (4070322.967491121 5643163.574441582)

Identifying STS in AIS data

A simple algorithm:

- **If** two or more ships
- **Spend** two or more hours
- **Within** 500m of each other
- **Then** link them together



SELECT

a1.vesselid AS ship1,

a2.vesselid AS ship2,

a1.date AS start,

a2.date AS end

FROM

ais AS a1

JOIN

ais AS a2

ON ST_DWithin(a1.geom, a2.geom, 500)

AND ABS(EXTRACT(EPOCH FROM (a2.date - a1.date))) > 7200;

AND a1.vesselid <> a2.vesselid

Spatial Joins

What neighborhood is the 'Broad St' station in?

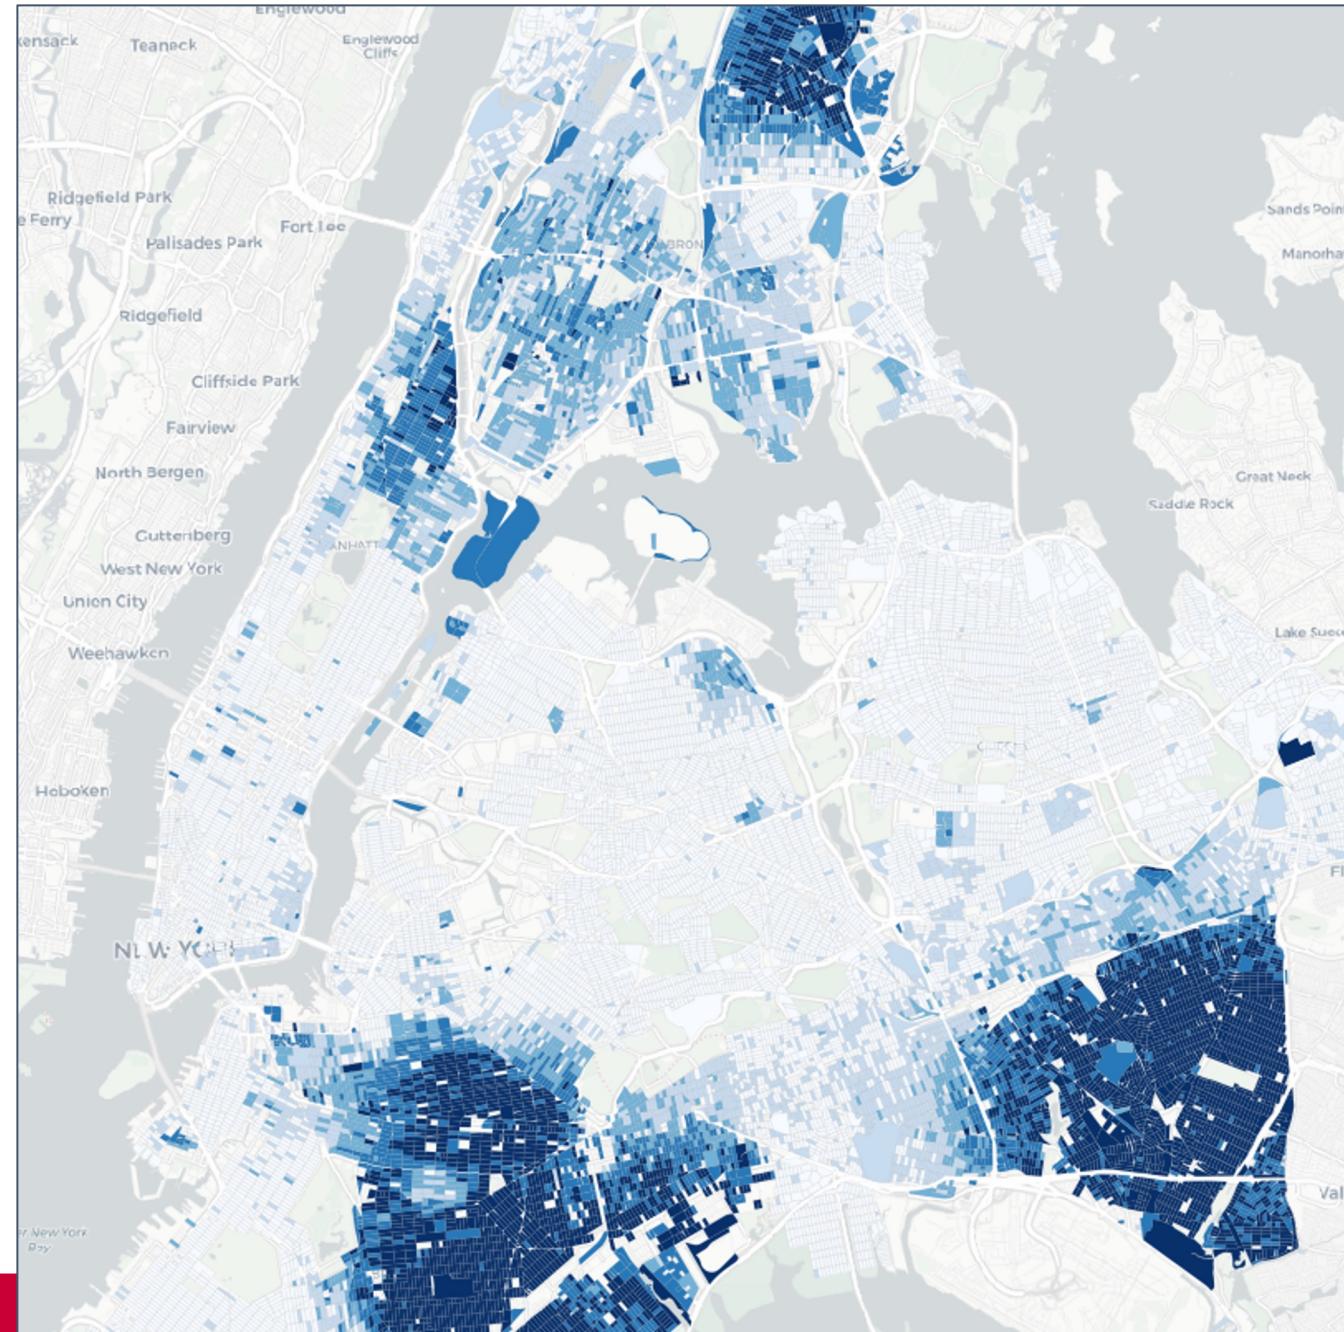
Remember...

```
SELECT name, boroname
FROM nyc_neighborhoods
WHERE
    ST_Intersects(
        geom,
        ST_GeomFromText(
            'POINT(583571 4506714)',
            26918));
```

Do it in one step, with a spatial join!

```
SELECT s.name, n.name, n.boroname
FROM nyc_neighborhoods AS n
JOIN nyc_subway_stations AS s
ON ST_Contains(
    n.geom,
    s.geom
)
WHERE s.name = 'Broad St';
```

What is the population and racial make-up of the neighborhoods of Manhattan?



```
SELECT
```

```
    n.name AS neighborhood_name,  
    SUM(c.popn_total) AS population,  
    100*SUM(c.popn_white)/SUM(c.popn_total) AS white_pct,  
    100*SUM(c.popn_black)/SUM(c.popn_total) AS black_pct
```

```
FROM nyc_neighborhoods AS n
```

```
JOIN nyc_census_blocks AS c
```

```
ON ST_Intersects(
```

```
    n.geom,
```

```
    c.geom
```

```
)
```

```
WHERE n.boroname = 'Manhattan'
```

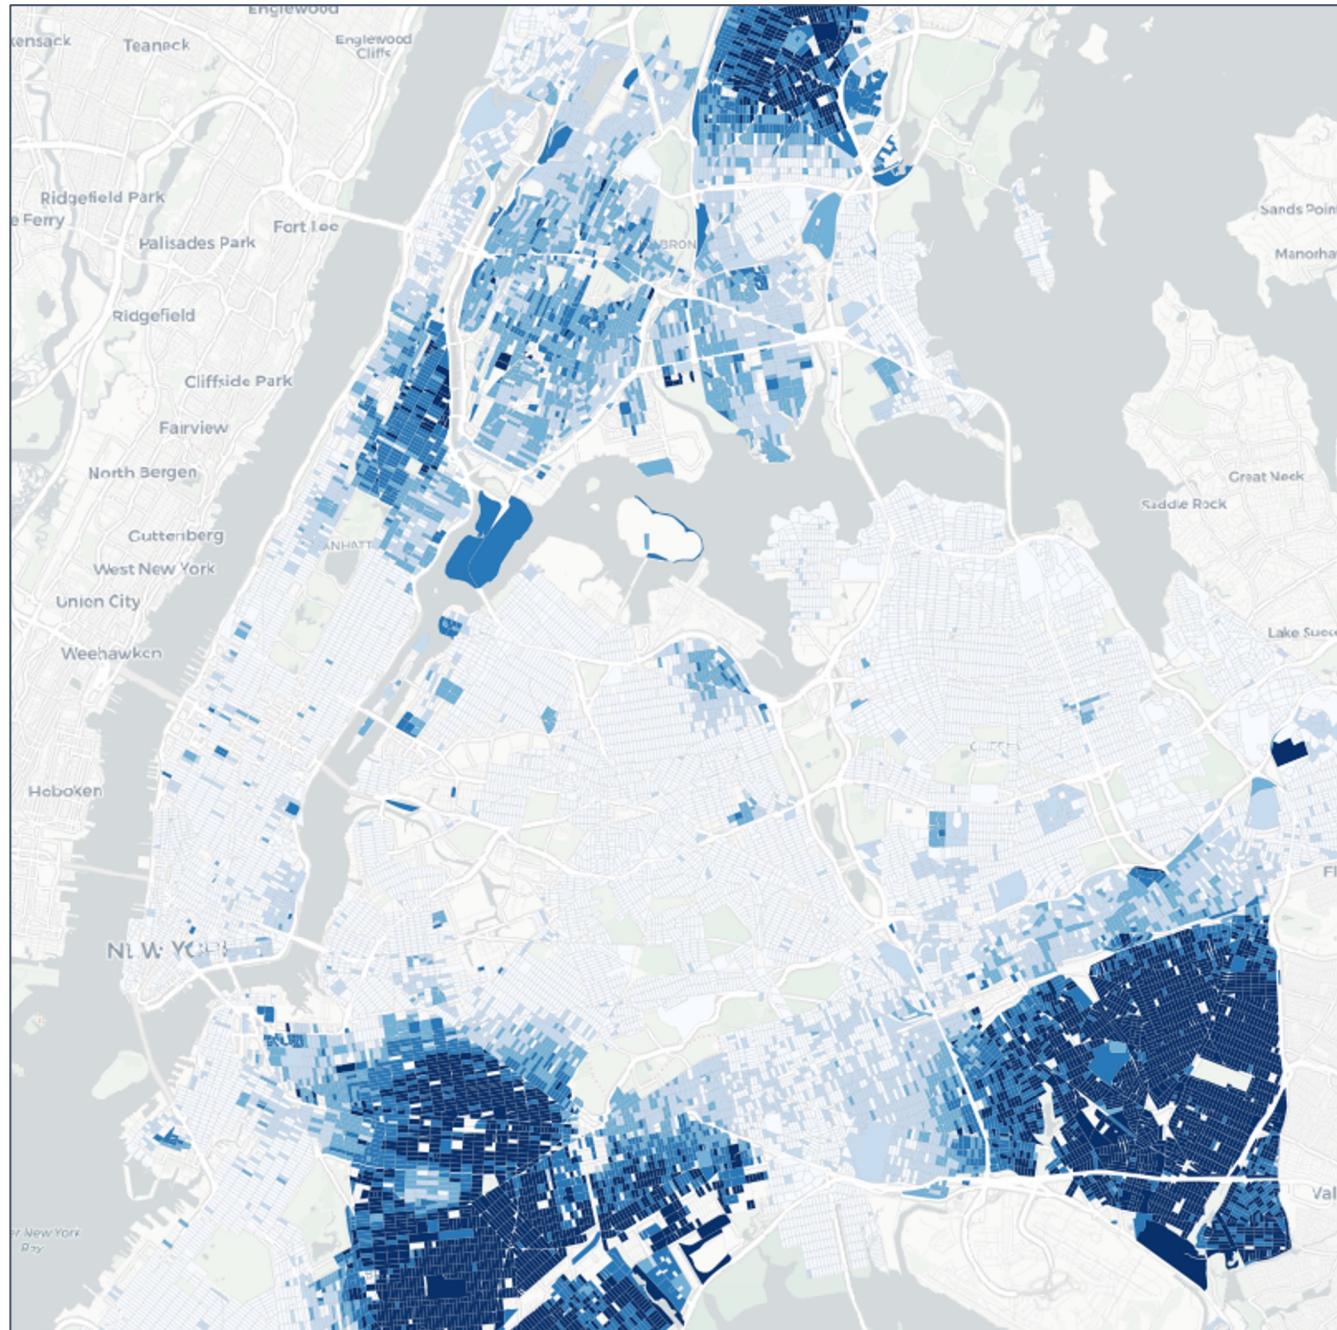
```
GROUP BY n.name
```

```
ORDER BY white_pct DESC;
```

neighborhood_name	popn	white %	black %
Carnegie Hill	18763	90.1	1.4
West Village	26718	87.6	2.2
North Sutton Area	22460	87.6	1.6
Upper East Side	203741	85.0	2.7
Soho	15436	84.6	2.2
Greenwich Village	57224	82.0	2.4
Central Park	46600	79.5	8.0
Tribeca	20908	79.1	3.5
Gramercy	104876	75.5	4.7
Murray Hill	29655	75.0	2.5
Chelsea	61340	74.8	6.4
Upper West Side	214761	74.6	9.2
Midtown	76840	72.6	5.2
Battery Park	17153	71.8	3.4

neighborhood_name	popn	white %	black %
Financial District	34807	69.9	3.8
Clinton	32201	65.3	7.9
East Village	82266	63.3	8.8
Garment District	10539	55.2	7.1
Morningside Heights	42844	52.7	19.4
Little Italy	12568	49.0	1.8
Yorkville	58450	35.6	29.7
Inwood	50047	35.2	16.8
Washington Heights	169013	34.9	16.8
Lower East Side	96156	33.5	9.1
East Harlem	60576	26.4	40.4
Hamilton Heights	67432	23.9	35.8
Chinatown	16209	15.2	3.8
Harlem	134955	15.1	67.1

**Let's explore the
racial geography
of New York City...**



Overall Racial Make-up of NYC

SELECT

```
100.0*SUM(popn_white)/SUM(popn_total) AS white_pct,  
100.0*SUM(popn_black)/SUM(popn_total) AS black_pct,  
SUM(popn_total) AS popn_total  
FROM nyc_census_blocks;
```

white_pct	black_pct	popn_total
44.00395007628105	25.546578900241613	8175032

**What is the racial make-up of
the areas served by the A train?**

First, we must determine where the A train stops.

Our routes are comma-separated strings!

```
SELECT DISTINCT routes  
FROM nyc_subway_stations;
```

4,5
N,Q,R,W
J
B,M,Q,R
D,F,N,Q
J,M
E,F
...

Postgres string function: **strpos()**

strpos(routes, 'A') returns a non-zero number if 'A' is in the routes field

Check out <https://www.postgresql.org/docs/current/functions-string.html>

Find all routes with an “A”

```
SELECT DISTINCT routes  
FROM nyc_subway_stations AS subways  
WHERE strpos(subways.routes, 'A') > 0;
```

A,C

A,B,C,D

A,C,E,L

A,C,F

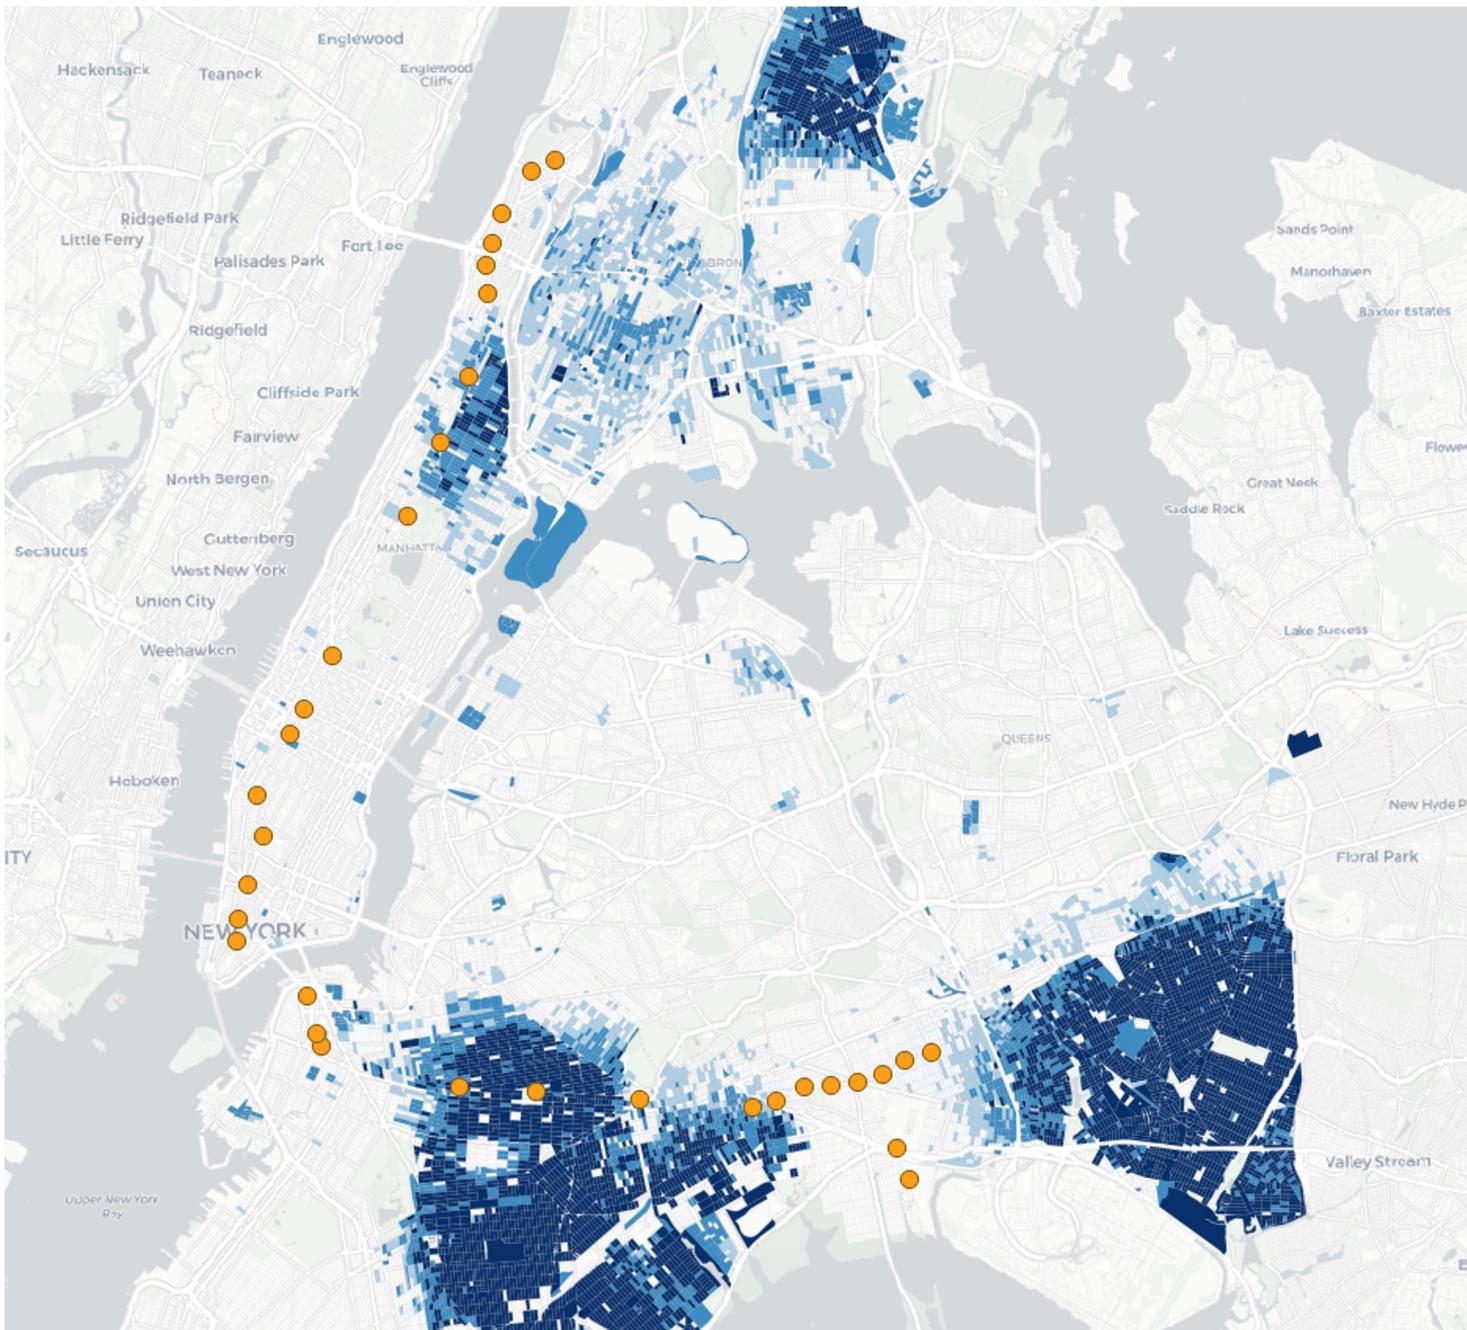
A,B,C

A,S

A,C,E

The route of the A train.

What is the racial makeup within 200 meters of each stop? Who is served by the A train?



Summarize population 200m from A train stops

SELECT

```
100*SUM(c.popn_white)/SUM(c.popn_total) AS white_pct,  
100*SUM(c.popn_black)/SUM(c.popn_total) AS black_pct,  
SUM(popn_total) AS popn_total  
FROM nyc_census_blocks AS c  
JOIN nyc_subway_stations AS s  
ON ST_DWithin(  
    c.geom,  
    s.geom,  
    200  
)  
WHERE strpos(s.routes, 'A') > 0;
```

New York

44.00% white

25.55% black

A Train

45.59% white

22.09% black

Spatial Indexing

A spatial database has...

- **Spatial Data Types**
 - geometry, geography
- **Spatial Indexes**
 - r-tree, quad-tree, kd-tree
- **Spatial Functions**
 - ST_Length(geometry), ST_X(geometry)
- NOTE: Currently, only PostGIS supports spatial indexing; DuckDB does not (but it's under development)!

A spatial index speeds spatial query ...

- Join two tables of 10,000 records each

Without Index	With Index
$10,000 * 10,000 = \mathbf{100,000,000}$ comparisons	$10,000 + 10,000 = \mathbf{20,000}$ comparisons

To prove it... remove the index.

```
DROP INDEX nyc_census_blocks_geom_idx;
```

Run a spatial join.

```
SELECT b.blkid  
FROM nyc_census_blocks b  
JOIN nyc_subway_stations s  
ON ST_Contains(b.geom, s.geom)  
WHERE s.name LIKE 'B%';
```

~300 ms

Create the index again.

```
CREATE INDEX nyc_census_blocks_geom_idx  
ON nyc_census_blocks USING GIST (geom);
```

Run the join again.

```
SELECT blocks.blkid  
FROM nyc_census_blocks b  
JOIN nyc_subway_stations s  
  ON ST_Contains(b.geom, s.geom)  
WHERE s.name LIKE 'B%';
```

~90 ms

Spatial Index Cliff Notes

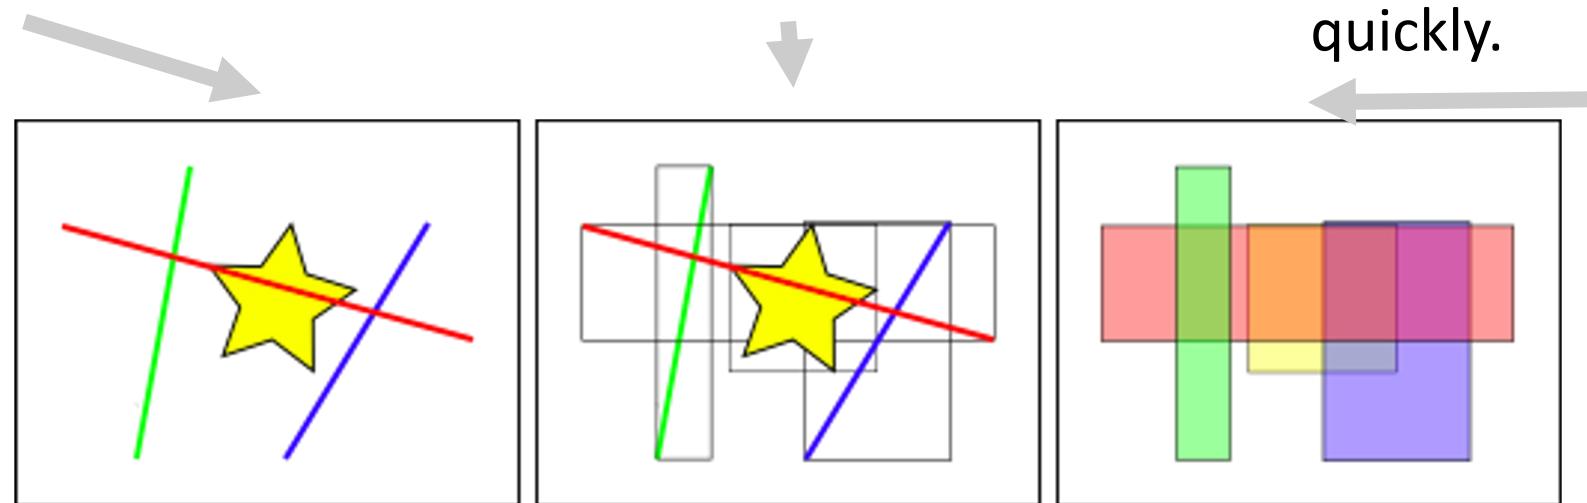
- **CREATE INDEX index_name ON table_name USING GIST (geom)**
- Use a “spatially indexed function” in **JOIN** or **WHERE** clause
 - **ST_Intersects(A, B), ST_Contains(A, B), ST_Within(A, B)**
 - **ST_DWithin(A, B, R)**

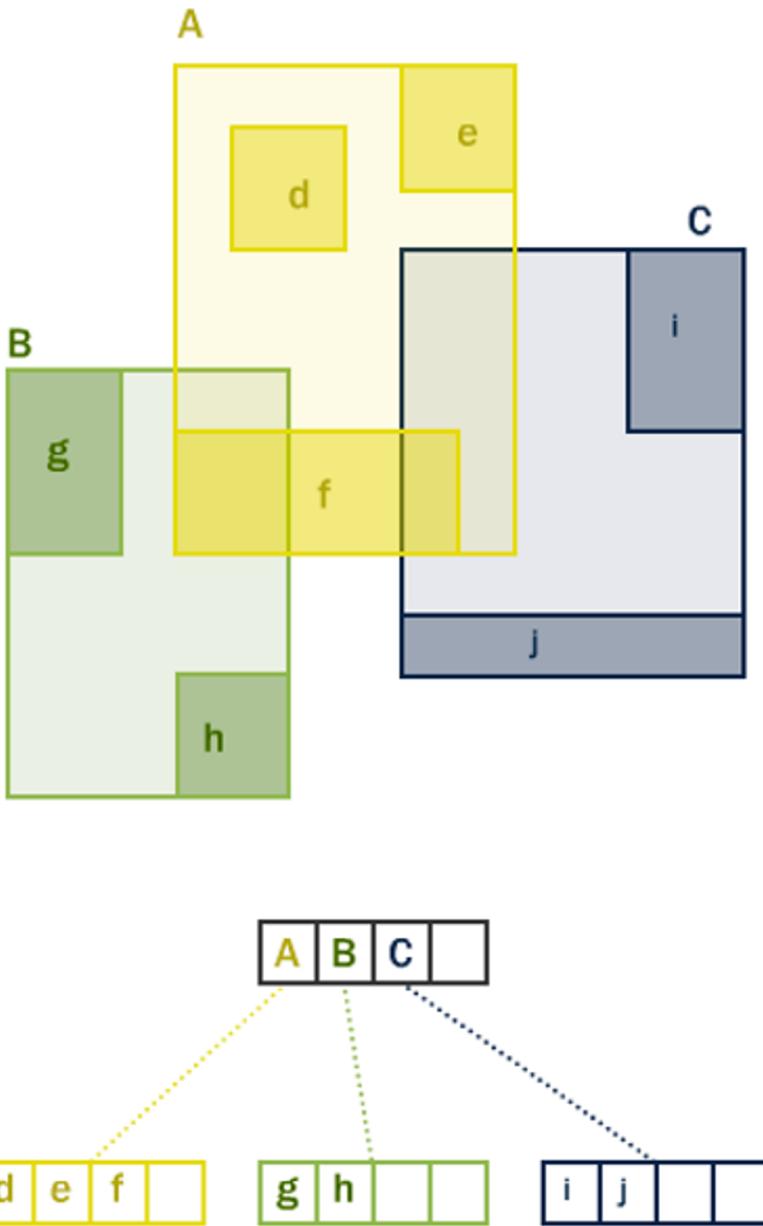
How do spatial indexes work?

Some spatial objects
(like the star) are quite
large and complex.
Comparing complex
objects is expensive!

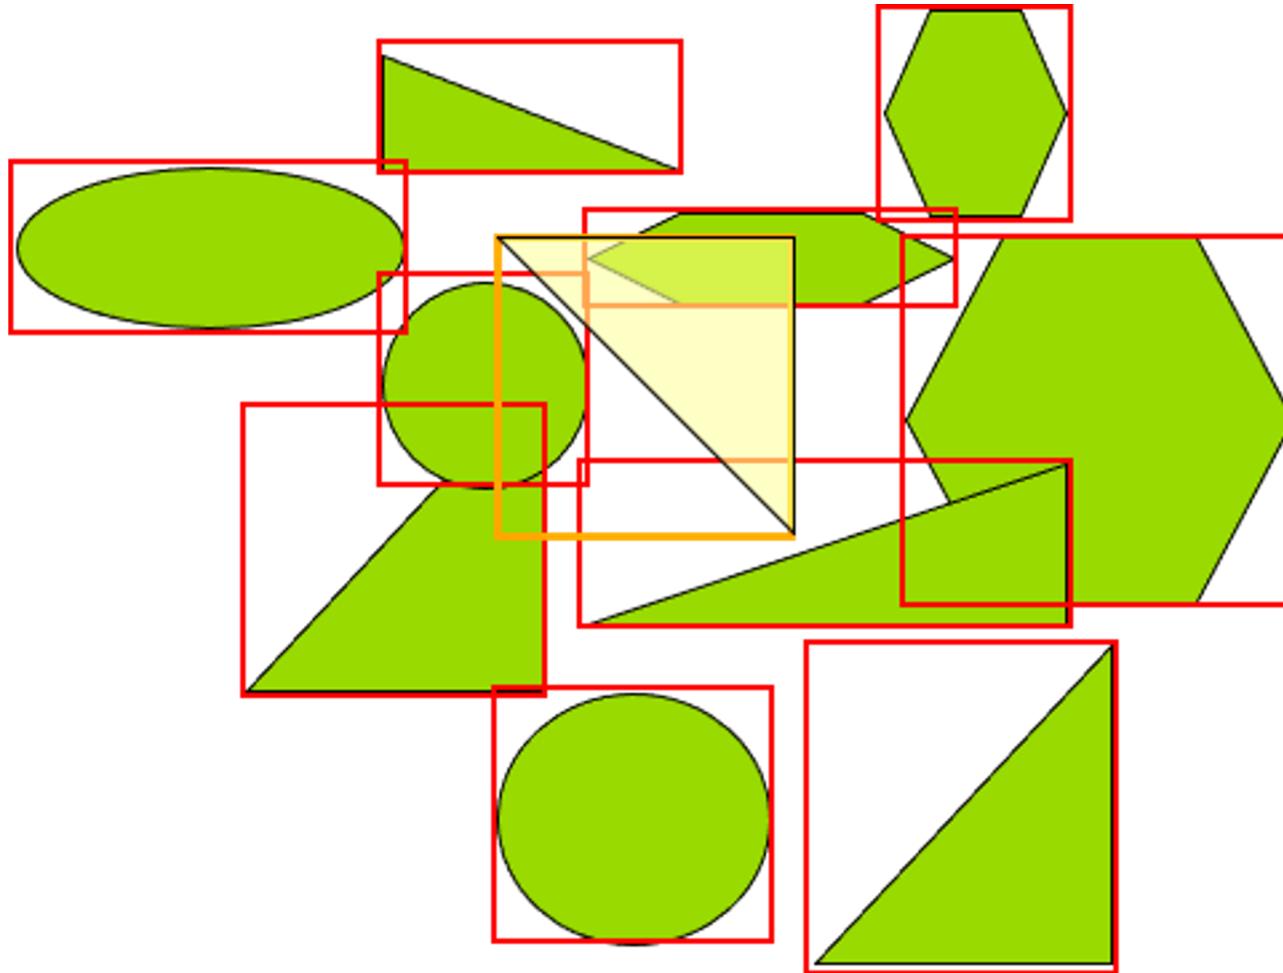
Instead of indexing
objects directly, spatial
indexes work on the
bounding boxes of the
objects.

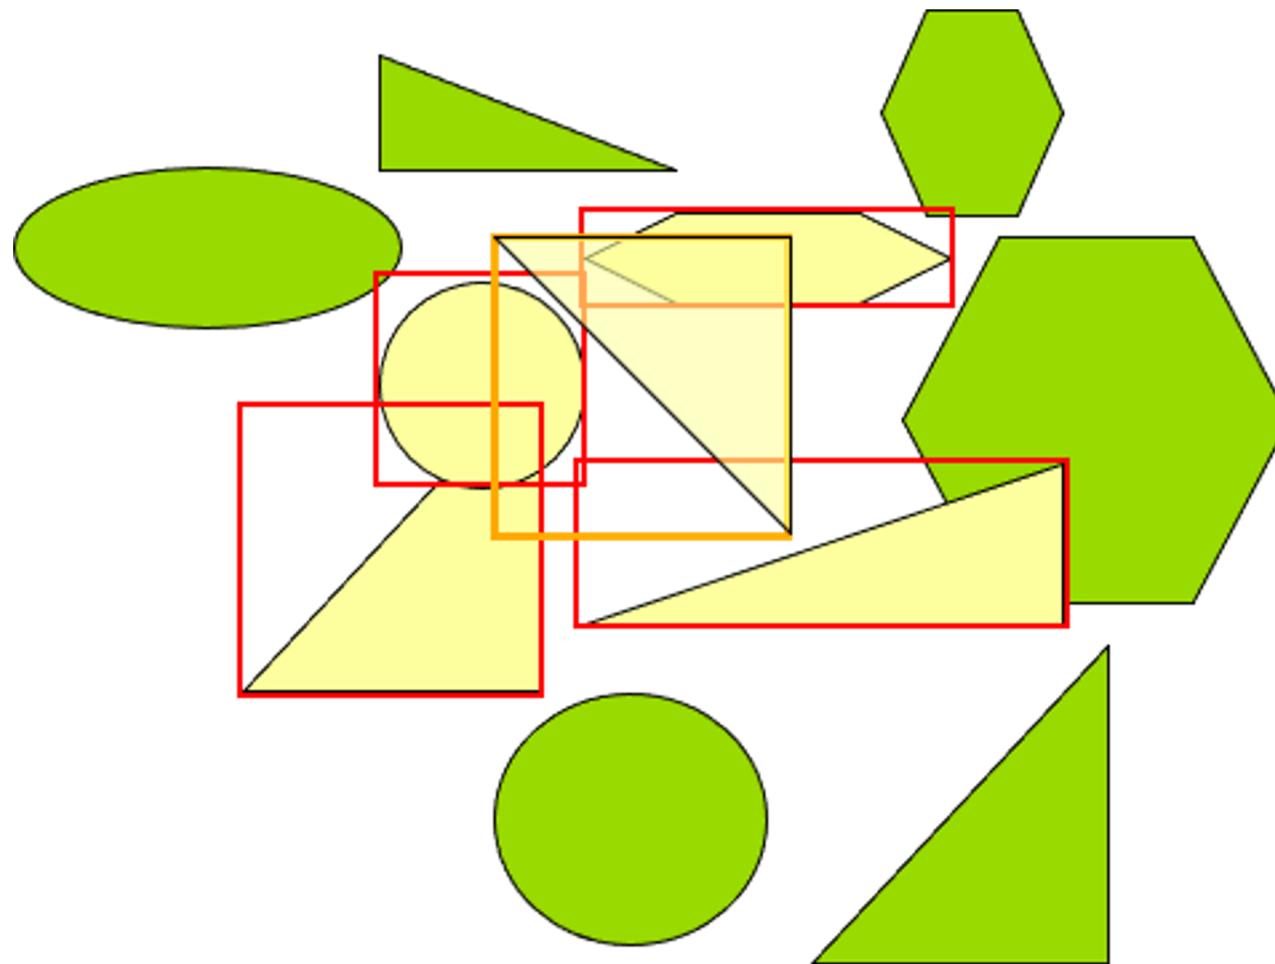
The boxes are of
uniform size, and can
be compared to
determine spatial
relationships very
quickly.

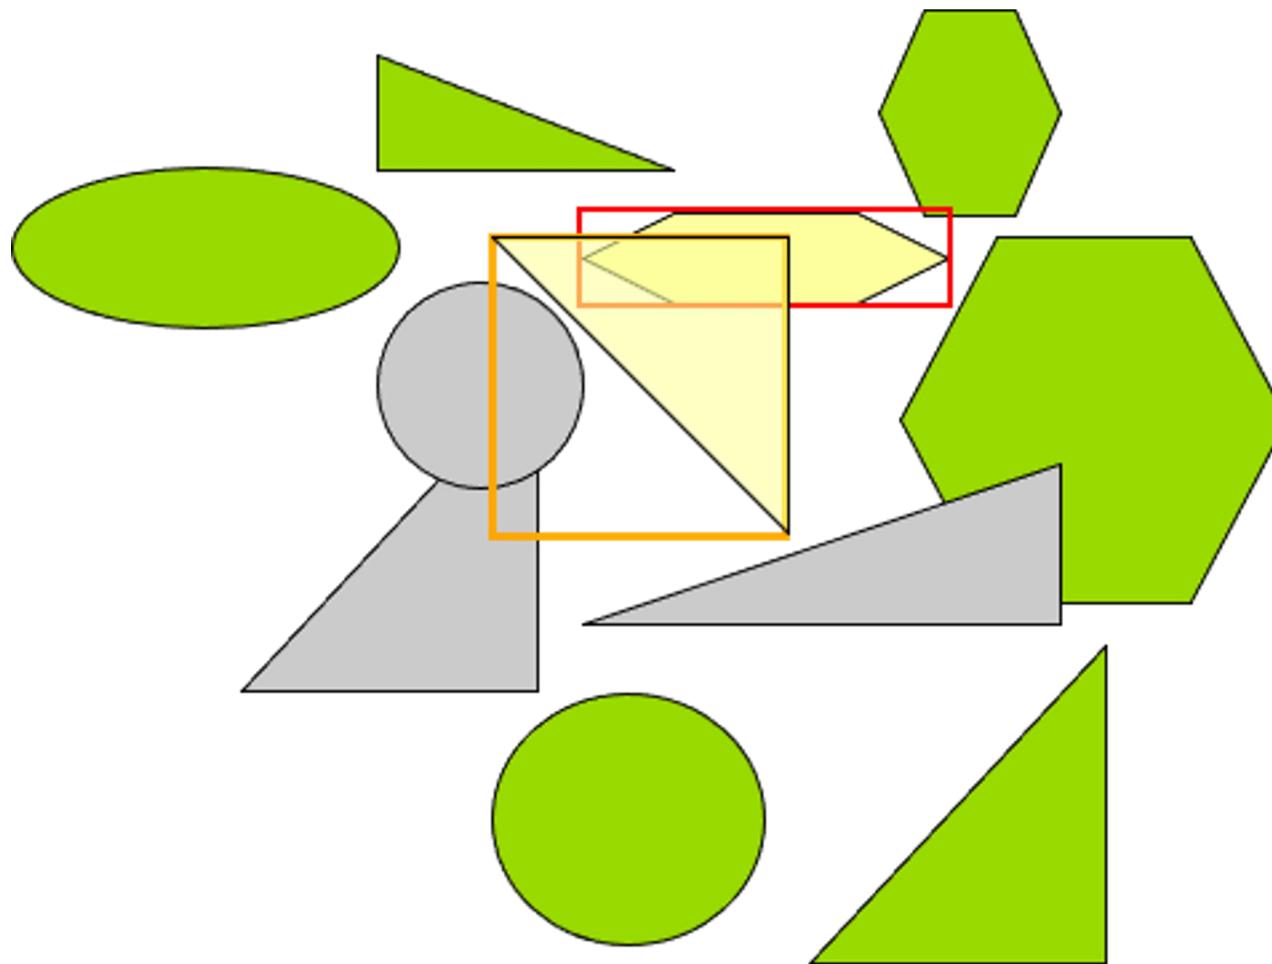




The boxes can be arranged in a hierarchy, so that a query can quickly discard portions of the search space that will not interact with a query box. Depending on the algorithm, different hierarchies can be build. PostGIS uses an “R*tree” algorithm.







Index-enabled Spatial Functions

- **ST_Intersects()**
- **ST_Contains()**
- **ST_Within()**
- **ST_DWithin()**
- ST_ContainsProperly()
- ST_CoveredBy()
- ST_Covers()
- ST_Overlaps()
- ST_Crosses()
- ST_DFullyWithin()
- ST_3DIntersects()
- ST_3DDWithin()
- ST_3DDFullyWithin()
- ST_LineCrossingDirection(
)
- ST_OrderingEquals()
- ST_Equals()

Index-only queries

`geom_a && geom_b`

The “`&&`” operator is the “bounding boxes overlap” operator.

It returns “true” when the bounds of the left and right arguments overlap.

Operators like “`=`” or “`>`” are symbols that express relationships between the left- and right-hand side arguments. “`&&`” is just another operator like any other.

What is the population of the West Village?

```
SELECT Sum(blk.popn_total)
FROM nyc_neighborhoods nh
JOIN nyc_census_blocks blk
ON nh.geom && blk.geom
WHERE nh.name = 'West Village';
```

49821

What is the population of the West Village?

```
SELECT Sum(blk.popn_total)
FROM nyc_neighborhoods nh
JOIN nyc_census_blocks blk
ON ST_Intersects(nh.geom, blk.geom)
WHERE nh.name = 'West Village';
```

26718

Spatial Indexes

```
VACUUM ANALYZE nyc_census_blocks;
```

The database planner can only know when to use indexes if the tables have been “analyzed”.

The database executor will run more efficiently if dead tuple bloat has been removed with “vacuum”. Most important on tables with bulk data changes (insert/update/delete).



Any Questions?

o.ballinger@ucl.ac.uk